

SemiticGPT: Efficient Multilingual Pretraining for a Script-Diverse, Low-Resource Language Cluster

Ronnen Slasky¹

¹Independent Researcher ronnen.slasky@gmail.com

Model: [Slasky/SemiticGPT-3B](#) · Code: [fatherRonnen/semitic-gpt](#)

April 2026

ABSTRACT

We present SemiticGPT, a 3.14-billion parameter decoder-only language model trained from scratch for Hebrew, Arabic, Persian (Farsi), and English—a script-diverse, low-resource language cluster centered on the Semitic family. The model is pretrained on 4.48 billion tokens using cloud spot instances at a total cost under \$1,500.

Our central finding is that **multilingual SFT is necessary and English-only SFT does not transfer cross-lingually** in our setup. Across a three-stage evaluation (base → English-only SFT → multilingual SFT), English-only fine-tuning produces near-zero improvement on non-English sentiment tasks, while multilingual SFT yields substantial gains: Hebrew 53%→84.5%, Arabic 45%→60.5%, Farsi 60.5%→78.5% (logprob classification on balanced evaluation sets).

We also demonstrate that a custom balanced tokenizer achieves **49–69% fewer tokens** than Llama-2 for Hebrew, Arabic, and Farsi at equivalent English efficiency, and that the model develops strong cross-lingual embeddings (90% retrieval accuracy for English↔Hebrew) purely from multilingual pretraining.

We provide a practitioner’s recipe for building similar models for other under-resourced language clusters. All code, data, and model weights are released.

Keywords: multilingual language models, cross-lingual transfer, Hebrew, Arabic, Farsi, Semitic languages, tokenizer design, low-resource NLP, efficient pretraining

1. Introduction

Large language models have achieved remarkable capabilities, yet their development remains concentrated on high-resource languages [Joshi et al., 2020]. Languages of the Semitic family—particularly Hebrew and Arabic—are underserved despite a combined speaker population exceeding 400 million. While Jais [Sengupta et al., 2023] addresses Arabic and various Hebrew models exist [Slasky, 2026], no prior work trains a generative model jointly covering Hebrew, Arabic, and Persian from scratch.

This paper addresses a specific empirical question: *Can efficient multilingual pretraining on a linguistically-motivated language cluster produce useful cross-lingual transfer, and what are the requirements for downstream task transfer?*

We investigate this through three experiments:

1. **Experiment A** (Hebrew Downstream): How well does multilingual SFT enable sentiment classification across all four languages?
2. **Experiment B** (Cross-lingual Transfer): Does English-only SFT transfer to non-English languages, or is multilingual SFT required?
3. **Experiment C** (Tokenizer Ablation): How much efficiency does a custom balanced tokenizer provide over English-centric alternatives?

Hebrew and Arabic share the Semitic triconsonantal root system, derivational morphology (binyan/wazn patterns), and

flexible word order. Persian, while using Arabic script and containing extensive Arabic loanwords, belongs to the Indo-European family—providing a control for distinguishing linguistic relatedness from script similarity.

Our primary contributions:

1. A **reproducible, low-cost recipe** (<\$1,500) for training a 3B multilingual model from scratch for an under-resourced language cluster.
2. Evidence that **English-only SFT does not transfer** cross-lingually at this scale—multilingual SFT is necessary for non-English task performance.
3. A custom tokenizer achieving **49–69% fewer tokens** than Llama-2 for target languages, with detailed ablation against two baselines.
4. A **practitioner’s guide** with step-by-step recommendations and negative results for researchers targeting other language clusters.

2. Related Work

Arabic Language Models. Jais [Sengupta et al., 2023] is a 13B bilingual Arabic-English model trained on 395B tokens. AraGPT2 [Antoun et al., 2021] provides GPT-2 scale Arabic models. ArabianGPT [Koubâa and Ammar, 2024] focuses on native Arabic without English contamination. None include Hebrew or Persian.

Hebrew-Arabic Models. HeArBERT [Rivas et al., 2024] explores bilingual Arabic-Hebrew modeling through transliteration-based script unification, but is limited to

⁰Code: <https://github.com/fatherRonnen/semitic-gpt>.
Model: <https://huggingface.co/Slasky/SemiticGPT-3B>.

encoder-only architecture. To our knowledge, no prior work trains a generative decoder-only model jointly on Hebrew, Arabic, and Persian from scratch.

Small Multilingual Models. SmolLM3 [SmolLM3, 2025] provides a 3B training recipe but targets European languages. The insights do not address non-Latin scripts or Semitic-family challenges.

Training Recipe Discovery. Our pretraining recipe was validated through proxy experiments at small scale following Karpathy [2026], as described in our prior work on Hebrew language modeling [Slasky, 2026]. We conducted 32 proxy experiments at 110M scale before committing to the 3B pre-training run.

3. Methodology

3.1 Tokenizer Design

We train a 32,000-token SentencePiece BPE tokenizer on a balanced sample from all four languages (approximately 25% each). This ensures no single language dominates the vocabulary.

Special tokens: `<|user|>`, `<|assistant|>`, `<s>` (BOS), `</s>` (EOS), `<pad>`.

We evaluate tokenizer quality in Experiment C (Section 4.4).

3.2 Pretraining Data

We collect approximately 4.48 billion tokens from publicly available sources with a Hebrew-anchored distribution:

- **Hebrew (40%):** Wikipedia, OSCAR, news corpora, government documents
- **Arabic (20%):** Wikipedia, OSCAR, news, UN parallel corpus
- **English (20%):** Wikipedia, OpenWebText subset
- **Persian (20%):** Wikipedia, OSCAR, news corpora




Figure 1: Pretraining data distribution. Hebrew is overrepresented as the anchor language.

Hebrew is intentionally overrepresented to ensure strong anchor-language representations. Data preprocessing includes MinHash deduplication, perplexity-based quality filtering, and language identification verification.

3.3 Architecture

SemiticGPT uses a decoder-only transformer with modern enhancements:

Hyperparameter	Value
Parameters	3.14B
Layers	26
Hidden dimension	3,072
Attention heads	24
Head dimension	128
Vocabulary size	32,000
Max sequence length	2,048
Position encoding	RoPE ($\theta=10,000$)
Activation	SwiGLU
Normalization	RMSNorm ($\epsilon=1e-6$)

Table 1: SemiticGPT architecture.

3.4 Pretraining

Training uses AWS spot instances (L40S 48GB and H100 80GB). Configuration:

- **Optimizer:** AdamW ($\beta_1=0.9$, $\beta_2=0.95$)
- **Learning rate:** $3e-4$ peak, cosine decay to 3%
- **Warmup:** 2,000 steps
- **Batch size:** 512K tokens (via gradient accumulation)
- **Weight decay:** 0.1 **Gradient clipping:** 1.0
- **Precision:** bfloat16

Total cost: approximately \$1,456 using spot instances at 60–70% discount.

3.5 Supervised Fine-Tuning (SFT)

We conduct SFT using 36,980 multilingual samples:

- **Sentiment data** (18,980 samples): Binary sentiment classification across all four languages, sourced from native-language datasets (Hebrew: HebEMO; Arabic: HARD; Persian: DeepSentiPers; English: SST-2).
- **Translation data** (18,000 samples): Parallel sentence pairs from OPUS-100 [Zhang et al., 2020], covering all six language-pair directions.

Data quality. After identifying critical data bugs in earlier SFT iterations (Section 5.5), the final dataset enforces: (1) correct label alignment verified per-dataset, (2) random shuffling before train/eval split, (3) strict separation with no data leakage, and (4) balanced 50/50 class distribution in all evaluation sets.

SFT training: AdamW optimizer, learning rate $2e-5$, cosine schedule, 8,000 steps, sequence length 384, bfloat16 precision. Best validation loss: 1.98.

For the cross-lingual transfer experiment (Experiment B), we also train an **English-only SFT** variant using only the English subset (4,745 samples) with identical hyperparameters for 2,000 steps.

4. Experiments and Results

We evaluate SemiticGPT through three focused experiments designed to answer specific questions about multilingual pre-training efficiency.

4.1 Evaluation Methodology

Sentiment classification. We use two complementary methods:

- **Logprob:** Compare log-probabilities of “positive” vs “negative” completion tokens given the input text and a classification prompt. This avoids generation parsing issues. $N=200$ per language.
- **Generative:** The model generates a free-form classification response. Outputs are parsed for sentiment labels. $N=50$ per language.

All evaluation sets are balanced (50% positive, 50% negative), so chance performance is 50% for logprob and 0% for generative (where the model must produce a parseable label).

Belebele [Bandarkar et al., 2023]: 4-way multiple-choice reading comprehension. $N=100$ per language. Chance = 25%.

Cross-lingual retrieval: Given a sentence in language A, retrieve its translation from 10 candidates in language B using cosine similarity of mean-pooled hidden states. Chance = 10%.

4.2 Experiment A: Multilingual Downstream Evaluation

We evaluate sentiment classification at three model stages: base (pretrained only), English-only SFT, and multilingual SFT.

Stage	HE	AR	FA	EN
<i>Logprob classification ($N=200$, chance=50%)</i>				
Base	53.0	45.0	60.5	51.5
EN-SFT	51.5	46.5	58.5	52.0
Multi-SFT	84.5	60.5	78.5	73.0
<i>Generative classification ($N=50$, chance=0%)</i>				
Base	0	0	0	0
EN-SFT	0	0	0	2
Multi-SFT	82	64	74	64

Table 2: Sentiment accuracy (%) across three model stages. Multilingual SFT produces large gains across all languages; English-only SFT does not.

Key observations:

1. **Base model:** Near chance for logprob (50–60%), zero for generative. The pretrained model has no task-following ability.
2. **English-only SFT:** Negligible improvement on any language including English itself (+0.5pp logprob). The model learns to follow English instructions but this does not transfer to sentiment understanding.
3. **Multilingual SFT:** Large gains across all languages. Hebrew improves most (+31.5pp logprob, 82% generative), consistent with its role as the pretraining anchor language.




Figure 2: Sentiment classification across three stages. English-only SFT (blue) produces no meaningful improvement over base (gray). Multilingual SFT (green) yields large gains for all languages.

4.3 Experiment B: Cross-lingual Transfer Analysis

Experiment B directly tests whether fine-tuning in one language transfers to others. Table 2 already shows the three-stage progression. The key finding is the **EN-SFT control**: training on English-only data does not improve non-English sentiment performance, and barely improves English itself.

This contrasts with findings in larger models where English fine-tuning can transfer cross-lingually [Conneau et al., 2020]. At 3B scale with 4.48B pretraining tokens, **explicit multilingual SFT data is required** for each target language.

Belebele reading comprehension (Table 3) shows a similar pattern: all stages perform near chance, confirming that 3B scale is insufficient for abstract reasoning tasks regardless of SFT strategy.

Stage	HE	AR	FA	EN
Random	25.0	25.0	25.0	25.0
Base	26.0	22.0	19.0	19.0
EN-SFT	26.0	23.0	20.0	20.0
Multi-SFT	28.0	24.0	25.0	23.0

Table 3: Belebele reading comprehension accuracy (% , $N=100$). All near chance (25%), as expected at 3B scale.

Cross-lingual embeddings. Despite limited task transfer, the base model develops strong cross-lingual alignment purely from multilingual pretraining:

Direction	Base	Multi-SFT
EN→HE	90%	90%
HE→EN	90%	80%
AR→EN	70%	50%
AR→HE	70%	70%
EN→AR	50%	60%
HE→AR	40%	60%
FA→HE	50%	30%
HE→FA	40%	30%
EN→FA	20%	20%
AR→FA	10%	10%
FA→EN	10%	10%
FA→AR	10%	10%
Average	45.8%	43.3%

Table 4: Cross-lingual retrieval accuracy (10-way, chance=10%). Strong EN↔HE alignment emerges from pretraining alone.

English↔Hebrew achieves 90% retrieval ($9\times$ chance), and Arabic shows meaningful alignment with both English and Hebrew. Persian shows weak connectivity, consistent with its typological distance from the Semitic family.

	EN	HE	AR	FA
EN	—	90%	70%	20%
HE	90%	—	40%	40%
AR	50%	70%	—	10%
FA	10%	50%	10%	—

Figure 3: Cross-lingual retrieval heatmap (base model, 10-way, chance=10%). Darker = stronger alignment. EN↔HE alignment is strongest.

4.4 Experiment C: Tokenizer Ablation

We compare our custom 32K tokenizer against two baselines of identical vocabulary size: Llama-2 (English-centric) and HebrewGPT-1B (Hebrew-centric). All three use 32K BPE vocabularies.

4.4.1 Vocabulary Composition

Script	Ours	Llama-2	HebrewGPT
Arabic script	46.7%	0.2%	0.4%
Hebrew script	27.8%	0.1%	72.2%
Latin script	24.3%	80.9%	7.0%
Other	1.2%	18.8%	20.4%

Table 5: Vocabulary composition by script. Our tokenizer balances Arabic (covering both Arabic and Persian), Hebrew, and Latin. Llama-2 is 81% Latin; HebrewGPT is 72% Hebrew.

4.4.2 Encoding Efficiency

Lang	Ours	Llama-2	HebrewGPT	Δ vs Llama
<i>Fertility (tokens per word, lower is better)</i>				
EN	1.54	1.51	2.42	−2.0%
HE	1.34	3.91	1.26	+65.7%
AR	2.22	4.36	4.15	+49.1%
FA	1.52	4.88	4.51	+68.8%
<i>Bytes per token (higher is better)</i>				
EN	3.71	3.79	2.37	−2.1%
HE	5.12	1.76	5.48	+191%
AR	3.48	1.77	1.86	+97%
FA	5.72	1.78	1.93	+221%

Table 6: Tokenizer efficiency comparison. Our tokenizer matches Llama-2 on English while using 49–69% fewer tokens for Hebrew, Arabic, and Persian. HebrewGPT excels on Hebrew but fragments Arabic/Persian.




Figure 4: Tokenizer fertility comparison. Lower is better. Our tokenizer achieves near-parity with Llama-2 on English while dramatically reducing fragmentation for Hebrew, Arabic, and Persian.

Key findings:

- No English penalty:** Our tokenizer matches Llama-2 on English (1.54 vs 1.51 tokens/word), a common concern when balancing non-Latin scripts.
- Massive efficiency for target languages:** Hebrew uses 66% fewer tokens, Arabic 49% fewer, Persian 69% fewer than Llama-2. This directly translates to longer effective context windows and reduced training cost.
- Specialization tradeoff:** HebrewGPT-1B slightly outperforms ours on Hebrew (1.26 vs 1.34) but fragments Arabic/Persian as badly as Llama-2. A multilingual tokenizer must balance all target languages.

4.4.3 Segmentation Examples

To illustrate the practical impact, consider the Arabic sentence “ .” (The capital of France is Paris):

- Ours:** 5 tokens (| | | | |)
- Llama-2:** 22 tokens (character-level fragmentation)
- HebrewGPT:** 19 tokens (character-level fragmentation)

Our tokenizer captures whole Arabic words as single tokens, while the baselines fragment at the character level—a 4.4× efficiency difference on this sentence.

5. Analysis and Discussion

5.1 The Multilingual SFT Requirement

The most important finding is the sharp contrast between English-only and multilingual SFT (Figure 2). At 3B scale with 4.48B pretraining tokens, the model does not develop sufficient cross-lingual task transfer from English fine-tuning alone. This differs from observations in much larger models (175B+ parameters) where English fine-tuning shows some cross-lingual generalization.

Why does this matter for practitioners? Researchers building models for low-resource languages cannot rely on cheap English-only SFT and hope for cross-lingual transfer. They must invest in multilingual SFT data for each target language—but the good news is that relatively small amounts suffice (our SFT uses only 36,980 total samples).

5.2 Representation vs. Task Transfer

An intriguing asymmetry: the base model develops strong cross-lingual *representations* (90% EN↔HE retrieval) but cannot leverage them for *tasks* without explicit SFT. Multilingual pretraining builds a shared semantic space, but task-specific fine-tuning is the key that unlocks it.

5.3 Hebrew as Anchor Language

Hebrew shows the strongest results across all experiments: highest sentiment accuracy (84.5%), strongest embedding alignment with English, and best tokenizer coverage. This is consistent with Hebrew receiving 40% of pretraining data. Whether anchor overrepresentation produces *disproportionate* returns (beyond the linear expectation from more data) remains an open question.

5.4 Persian Isolation

Persian consistently shows weak cross-lingual connectivity despite sharing Arabic script. In our retrieval experiment, Persian achieves only 10–30% accuracy with other languages, compared to 40–90% for the Hebrew-Arabic-English triad. This suggests that **in our setup, linguistic family membership predicts transfer better than script similarity**, though we cannot fully disentangle this from Persian’s 20% pretraining share vs Hebrew’s 40%.

5.5 Data Quality Lessons

During development, we discovered critical data bugs that invalidated earlier results:

1. **Hebrew label inversion:** The dataset’s label encoding (0=positive) was opposite to our assumption (0=negative), causing the model to learn inverted sentiment.
2. **Arabic sorting artifact:** The dataset was sorted by label before splitting, causing the training set to contain only negative examples.

3. **Translation data quality:** Early iterations used 50 hard-coded sentence pairs instead of real parallel data.

Lesson: Automated data pipeline audits—checking label distributions, train/eval overlap, and sample diversity—should be mandatory before any training run. We now enforce 50/50 class balance and verify shuffle randomness in all evaluation sets.

5.6 Cost Analysis

Component	GPU-hours	Cost
Proxy experiments (32 runs)	16	\$12
Pretraining (4.48B tokens)	380	\$1,380
SFT (36,980 samples)	4	\$14
Evaluation (all experiments)	8	\$30
Total	408	~\$1,436

Table 7: Cost breakdown. All compute on AWS spot instances (60–70% discount).

6. Practitioner’s Guide

Based on our experience, we provide recommendations for researchers building similar multilingual models.

6.1 Language Cluster Selection

- Prioritize **linguistic family membership** over geographic proximity or script similarity. Our results suggest that Semitic family languages (Hebrew, Arabic) show stronger mutual transfer than script-mates (Arabic, Persian).
- Include one high-resource language (English) for evaluation infrastructure and instruction-following bootstrapping.
- Include a typologically distant language as a control condition.

6.2 Tokenizer Design

- Train a **custom BPE tokenizer** on equally-weighted samples of all target languages. Do *not* reuse an English-centric tokenizer—it wastes 49–69% of tokens on non-Latin scripts.
- **32K vocabulary** is sufficient for 4–6 languages. Target fertility ≤ 2.0 tokens/word for all languages.
- The English efficiency penalty from balancing is minimal (–2% in our case).

6.3 Pretraining

- **Anchor overrepresentation** (30–40% for primary language) appears beneficial but is not ablated in this work.
- Use **spot instances with checkpointing** every 30 minutes for cost efficiency.
- **bfloat16** is sufficient—we observed no degradation compared to float32 in preliminary experiments.
- Validate recipe at proxy scale (100–200M parameters) before full commitment.

6.4 Supervised Fine-Tuning

- **Include all target languages** in SFT data. English-only SFT does not transfer at 3B scale.
- **Native-language data > machine-translated data.** We found translated Alpaca produces weaker results than native-language datasets.
- **Audit data pipelines** before training: verify label encoding, check for sorting artifacts, confirm train/eval separation.
- 37K samples with 8K steps was sufficient for our model.

6.5 Evaluation

- Use **logprob-based evaluation** for classification tasks—it avoids prompt/parsing issues that plague generative evaluation.
- **Always report chance level** for each metric.
- Cross-lingual retrieval requires **no labeled data** and reveals representational quality.
- Expect **near-chance results on reasoning benchmarks** (Belebele, MMLU) at 3B scale. This is not a failure—it calibrates expectations.

Decision	Recommendation
Language selection	Linguistic family > script similarity
Tokenizer	Custom BPE, 32K vocab, balanced
Anchor language	Over-represent at 30–40%
Recipe validation	20–50 proxy experiments first
Optimizer	AdamW ($\beta_2=0.95$)
Precision	bfloat16
SFT data	All target languages required
SFT source	Native-language > translated
Data quality	Audit labels, splits, balance
Evaluation	Logprob > generative for classification
Compute	Cloud spot instances + checkpointing

Table 8: Recipe summary for building multilingual models from scratch.

7. Limitations

- **Single-run results:** All experiments are single-run without confidence intervals. Effects under 5 percentage points may not be statistically significant. Key findings (e.g., 84.5% vs 51.5% sentiment) are large enough to be meaningful.
- **Small evaluation sets:** 200 samples for logprob, 50 for generative, 100 for Belebele. Sufficient for detecting large effects but not small ones.
- **No external baselines:** We do not compare against mBERT, XLM-R, or NLLB on identical tasks. Such comparisons would contextualize our results.
- **Scale:** At 3B parameters, the model is below thresholds for complex reasoning. Findings about transfer requirements may not hold at larger scales where English-only SFT *does* transfer.
- **Confounded variables:** Hebrew receives 40% of pre-training data vs Persian’s 20%. Persian’s weak transfer may partly reflect data quantity rather than purely typological distance.

- **Anchor ablation:** We do not ablate the effect of anchor overrepresentation (e.g., 40% Hebrew vs 25% balanced).
- **Generalizability:** Our recipe is demonstrated on one language cluster. Whether the same principles apply to other families (Turkic, Bantu, Dravidian) is an empirical question.

8. Conclusion

We present SemiticGPT, a 3.14B parameter multilingual model covering Hebrew, Arabic, Persian, and English, trained from scratch for under \$1,500. Our three experiments establish that:

1. **Multilingual SFT is required:** English-only fine-tuning does not transfer to non-English tasks at 3B scale. Explicit multilingual SFT data unlocks the cross-lingual representations built during pretraining.
2. **Custom tokenizers eliminate waste:** A balanced tokenizer uses 49–69% fewer tokens than English-centric alternatives for target languages, with negligible English penalty.
3. **Representations precede tasks:** The model builds strong cross-lingual embeddings (90% EN↔HE retrieval) from pretraining alone, but task-specific SFT is needed to leverage them.

We release all code, data, and model weights to support reproducible research and provide a practitioner’s guide for researchers targeting other under-resourced language families.

Future work. Scaling to 7B+ parameters to test whether cross-lingual SFT transfer emerges at larger scale; adding more Semitic languages (Amharic, Tigrinya); ablating anchor language overrepresentation; and improving Arabic SFT with higher-quality native datasets.

9. Reproducibility

All artifacts are available at <https://github.com/fatherRonnen/semitic-gpt>:

Section	Repo Path	Contents
Tokenizer (§3.1)	<code>tokenizer/</code>	Config, training script, fertility
Data (§3.2)	<code>data/</code>	Collection scripts, manifests
Proxy (§2)	<code>proxy/</code>	32 configs, results
Pretraining (§3.4)	<code>pretrain/</code>	Scripts, hyperparams, logs
SFT (§3.5)	<code>sft/</code>	Configs, scripts, v4 pipeline
Experiments (§4)	<code>eval/</code>	Scripts, prompts, results

Table 9: Paper-to-repository mapping.

Model weights (base: 6.1GB, SFT: 6.3GB) and tokenizer are available on HuggingFace at <https://huggingface.co/Slasky/SemiticGPT-3B>.

References

W. Antoun, F. Baly, and H. Hajj. AraGPT2: Pre-trained transformer for Arabic language generation. In *Proc. WANLP*, 2021.

- L. Bandarkar et al. The Belebele benchmark: a parallel reading comprehension dataset in 122 language variants. *arXiv:2308.16884*, 2023.
- A. Conneau, K. Khandelwal, N. Goyal, et al. Unsupervised cross-lingual representation learning at scale. In *Proc. ACL*, 2020.
- P. Joshi, S. Santy, A. Buber, et al. The state and fate of linguistic diversity and inclusion in the NLP world. In *Proc. ACL*, 2020.
- A. Koubâa and A. Ammar. ArabianGPT: Native Arabic GPT-based large language model. *arXiv:2402.15313*, 2024.
- Marco-LLM Team. Marco-LLM: Bridging languages via massive multilingual training for cross-lingual enhancement. *arXiv:2412.04003*, 2024.
- A. Rivas et al. HeArBERT: A bilingual model for Arabic-Hebrew translation using transliteration. *arXiv:2402.16065*, 2024.
- N. Sengupta et al. Jais and Jais-chat: Arabic-centric foundation and instruction-tuned open generative large language models. *arXiv:2308.16149*, 2023.
- L.B. Allal et al. SmolLM3: The complete blueprint of a state-of-the-art 3B parameter LLM. HuggingFace Technical Report, 2025.
- A. Karpathy. autoresearch: Autonomous AI-driven ML experimentation. GitHub repository, 2026.
- R. Slasky. Autonomous AI-driven Hebrew language model optimization: From random search to optimizer-architecture co-discovery. Independent research report, 2026.
- B. Zhang, P. Williams, I. Titov, and R. Sennrich. Improving massively multilingual neural machine translation and zero-shot translation. In *Proc. ACL*, 2020.